

RecallRoom

A supplier recalls one ingredient lot.
Which products and customers did it reach?

An evidence workspace for small food manufacturers


recallroom.web.app

Shivam Gupta
Project creator

Nebius × NVIDIA
Best Apps and Agents

The recalled lot has a traceable scope

Fictional drill: peanut-butter lot PB-0901-A



Source records support each relationship. Finished units and ingredient kilograms stay separate.

A separate 240-cookie batch remains unresolved and on hold in the case.

Uncertainty changes only with evidence

Before review

240 cookies

The lot reference is ambiguous.
120 shipped units require investigation.
The case keeps this exposure visible.

40-50 kg

Lot A ingredient remaining

After a reviewed correction


Lot B confirmed

A reviewer attaches supporting evidence.
The decision log records the correction.
The graph recalculates Lot A exposure.

50 kg

Lot A ingredient remaining

The model drafts. The reviewer decides.



Firestore hosts the app. Cloud Run and Cloud Tasks handle authenticated extraction.

HTML, CSV and JSON exports include evidence and unsent customer drafts.

Hosted Nemotron 3 Super run: 17.562 seconds, about \$0.00456 in estimated inference cost.

*Text PDFs require a person to verify the extracted text. Scanned-image OCR is outside this MVP.

A focused workflow for a specific buyer

Small food manufacturers and co-packers

\$99

per site / month

Pricing hypothesis

No paying customers yet

Quality managers run a traceability drill using records they already keep. Consultants could bring the workflow to several client sites.

FDA's 2026 exercise with 15 voluntary firms found gaps in traceability information. The sample is not industry-wide evidence.

FoodDocs and Mar-Kov already offer traceability. RecallRoom's proposed entry point is a narrow incident workspace over existing documents.

Application evidence and next validation

Verified application checks

41 automated tests
60 hosted API checks
10 operator checks
Hosted NVIDIA flow passed

Next evidence

Evaluate more difficult source files
Measure against human answer keys
Observe quality managers using it
Test a paid pilot and repeat usage

One fictional scenario tested end to end. Real-world accuracy remains unmeasured.

Shivam Gupta

Project creator, with AI-assisted implementation