

RecallRoom

A reviewable response to an ingredient recall

Product brief / Created by Shivam Gupta

The problem and the buyer

A supplier names one ingredient lot. A small food manufacturer must work out which batches and customer shipments it reached. The relevant evidence may sit in receiving sheets, production records and shipment exports. RecallRoom gives the quality manager a private case to reconcile those records and expose what still needs review.

The workflow

Upload records and retain the originals. NVIDIA Nemotron through Nebius Token Factory proposes structured relationships with source references. Application checks validate the draft and calculate lot scope deterministically. A reviewer approves the draft, investigates ambiguity and exports the response packet.

The model interprets records. It cannot approve a recall decision or silently release a hold. Customer messages remain drafts that the application never sends.

A fictional drill with a consequential ambiguity

Lot PB-0901-A	Synthetic result
Confirmed finished products	480 oat bars + 320 energy bites = 800 units
Confirmed disposition	600 shipped to 3 customers; 200 on site
Separate cookie uncertainty	240 units, including 120 shipped
Lot A ingredient remaining	40-50 kg before review; 50 kg after correction

The cookie shorthand could name either received peanut-butter lot. A reviewed supporting record assigns it to lot B. The decision log preserves that change. Ingredient kilograms and finished-unit counts remain separate. These are demonstration facts, not customer outcomes or a general safety clearance.

The commercial hypothesis

Start with recurring traceability drills for small manufacturers and co-packers. Test **\$99 per site per month**, with food-safety consultants as a possible distribution channel. The price, recurring use and willingness to pay remain hypotheses. No customer interviews, paid pilots or measured savings have been completed.

FoodDocs and Mar-Kov already provide traceability capabilities. RecallRoom proposes a narrower incident workspace over existing exports and documents. Potential defensibility would come from reliable document adapters, reviewed lot aliases and a difficult evaluation set. Onboarding and support costs may determine viability more than inference cost.

Why this workflow deserves testing

FDA's 2026 tabletop exercise involved 15 voluntary companies. It found gaps in information retention and sharing, while many relevant facts existed in traditional business records. The sample does not represent the whole industry. RecallRoom targets evidence quality and readiness; it does not claim automatic FDA compliance. [1]

Implemented MVP and current proof boundary

Private authentication and saved cases in D1. Source originals in R2. Text, CSV, JSON and text-PDF uploads, with human verification of PDF text. Structured Nemotron drafts with human approval. Deterministic lot graphs and quantity ranges. An evidence decision log. HTML, CSV and JSON response packets with customer drafts.

Live Nebius execution is pending a credential. No measured live latency, model accuracy or inference cost is available yet. Local application checks and synthetic examples do not prove real-world reliability. Scanned-image OCR is outside the MVP.

The next validation

Record a successful live NVIDIA/Nebius run. Compare scope and provenance against a human-authored answer key, including ambiguity and duplicate imports. Observe quality managers completing a drill. Measure review effort and unsupported exclusions. Test an actual paid pilot and whether users repeat the workflow.

Primary sources and project

[1] FDA Traceability Readiness Tabletop Exercises Final Report (2026)

[2] FoodDocs: product and pricing

[3] Mar-Kov: manufacturing traceability

Public project repository

Created by Shivam Gupta, with AI-assisted research, implementation, testing and documentation.